



Recommender systems for clothing, shoes, and jewelry using Amazon data

User-based, item-based, and user/content hybrid methods
of collaborative filtering

George Nakhleh & Gunjan Pandya

Abstract:

Amazon is arguably the company that brought recommender systems to the masses. “Users who like ____ also like ____” is a familiar phrase to anyone who has made purchases from the website, and recommender systems are a cornerstone of its business. For analysts interested in getting hands-on experience with recommender systems, then, Amazon is a familiar reference point. Fortunately, datasets of Amazon product reviews are available online for just this reason.

This project uses a subset of reviews for a specific set of Amazon products (Clothing, Shoes & Jewelry) to design different types of recommender systems. The end result is three separate methods of delivering product recommendations: user-based, item-based, and a user/content hybrid that leverages product metadata. These recommenders were evaluated using mean absolute error and recall.

[Topics: recommender systems, collaborative filtering, user-based, item-based, content-based, cosine similarity, mean absolute error, recall]

About the data

A brief summary of the dataset:

Clothing, Shoes & Jewelry: subset of Amazon clothing, shoes, and jewelry products that ensures each product has at least 5 reviews and each reviewer has reviewed at least 5 products.

Format: JSON, converted to csv for quicker input.

Source: University of California, San Diego (UCSD) ([link](#))

Description: rows are reviews, variables described below:

asin: unique identifier for product, essentially a product id

helpful: List of two items - # of users who found review helpful ([2,5] = 2 of 5 users found this review helpful)

overall: user review of the product on a 5-star scale

reviewText: text of the review

reviewTime: date of the review (mm.dd.yyyy)

reviewerID: ID of the user

reviewerName: user name

summary: summary of the review text

unixReviewTime: time in unix format

Size: 278,677 rows, 9 columns (112 MB)

Snapshot of the dataset:

	asin	helpful	overall	reviewText	reviewTime	reviewerID	reviewerName	summary	unixReviewTime
0	0000031887	[0, 0]	5	This is a great tutu and at a really great pri...	02 12, 2011	A1KLRMWW2FWPL4	Amazon Customer "cameramom"	Great tutu- not cheaply made	1297468800
1	0000031887	[0, 0]	5	I bought this for my 4 yr old daughter for dan...	01 19, 2013	A2G5TCU2WDFZ65	Amazon Customer	Very Cutell	1358553600
2	0000031887	[0, 0]	5	What can I say... my daughters have it in oran...	01 4, 2013	A1RLQXYNCMWRWN	Carola	I have buy more than one	1357257600
3	0000031887	[0, 0]	5	We bought several tutus at once, and they are ...	04 27, 2014	A8U3FAMSJVHS5	Caromcg	Adorable, Sturdy	1398556800
4	0000031887	[0, 0]	5	Thank you Halo Heaven great product for Little...	03 15, 2014	A3GEOILWLK86XM	CJ	Grammy's Angels Love it	1394841600

Metadata: additional information about each clothing, shoes, and jewelry product.

Description: rows are products, variables described below

price: price of the item

salesRank: where the product ranks in a given sales category

categories: product categories that the given product falls under (list of lists)

imUrl: link to an image of the product

related: json-like data (dictionary) of similar items (eg 'also viewed', 'also bought')

asin: product id

title: name of the product

brand: brand of the product

description: description of the product

Size: 1,503,384 rows, 9 columns (1.36 GB)

Snapshot of the dataset:

	price	salesRank	categories	imUrl	related	asin	title	brand	description
0	6.99	{'Clothing': 1233557}	[[{'Clothing, Shoes & Jewelry', 'Girls'}, {'Clo...	http://ecx.images-amazon.com/images/I/31mCncNu...	{'also_viewed': ['B00JO8II76', 'B00DGN4R1Q', '...	0000037214	Purple Sequin Tiny Dancer Tutu Ballet Dance Fa...	Big Dreams	NaN
1	6.79	{'Sports & Outdoors': 8547}	[[{'Clothing, Shoes & Jewelry', 'Girls', 'Cloth...	http://ecx.images-amazon.com/images/I/314qZjYe...	{'also_bought': ['0000031852', '0000031895', '...	0000031887	Ballet Dress-Up Fairy Tutu	Boutique Cutie	This adorable basic ballerina tutu is perfect ...
2	64.98	{'Kitchen & Dining': 16987}	[[{'Clothing, Shoes & Jewelry', 'Novelty, Costu...	http://ecx.images-amazon.com/images/I/413tGhqo...	{'also_bought': ['B000BMTCK6', 'B0006JCGUM', '...	0123456479	SHINING IMAGE HUGE PINK LEATHER JEWELRY BOX / ...	NaN	Elegance par excellence. Hand-crafted of the f...
3	NaN	{'Clothing': 1180499}	[[{'Clothing, Shoes & Jewelry', 'Women', 'Acces...	http://ecx.images-amazon.com/images/I/31QZTHxv...	{'also_viewed': ['B008MTRT1O', 'B00BUG47S4', '...	0456844570	RiZ Women's Beautify Crafted ½ Rrimmed F...	NaN	NaN

Exploratory analysis

Before setting up recommender systems and their prerequisite components, some time was spent sizing up the initial dataset. Below are some details of the data worth noting.

Average number of user reviews: The average user has reviewed 7 products (median of 6)

Average number of product reviews: The average product has 12 reviews (median: 8)

Avg product rating: 4.245, so reviewers are generally pretty generous.

Preprocessing

First, a user-item matrix needed to be generated, an array where rows are users and columns are each unique product, with values being the users' reviews of a given product. This way, we can see who reviewed what products how.

The user-item matrix is used to determine similarity between users and between products: users who review the same products the same way are considered similar, and products reviewed by the same users the same way are too. The ability to find most similar users (or products) to a given user (or product) is the basis of collaborative filtering.

The code below creates a user-item matrix, replacing null values with 0.

```
data_df_slim = data_df[['asin', 'overall', 'reviewerID']]
user_item_mat = data_df_slim.pivot(index='reviewerID', columns = 'asin', values='overall')
test = np.array(user_item_mat)
nonan_test = np.nan_to_num(test)
```

The resulting matrix has 39,387 rows and 23,033 columns. It is very sparse: Of the 23,033 possible products, keep in mind that the average user has reviewed around 7.

Snapshot of the user-item matrix, an array

```
In [24]: nonan_test
Out[24]: array([[ 0.,  0.,  0., ...,  0.,  0.,  0.],
 [ 0.,  0.,  0., ...,  0.,  0.,  0.],
 [ 0.,  0.,  0., ...,  0.,  0.,  0.],
 ...,
 [ 0.,  0.,  0., ...,  0.,  0.,  0.],
 [ 0.,  0.,  0., ...,  0.,  0.,  0.],
 [ 0.,  0.,  0., ...,  0.,  0.,  0.]])
```

We can see above that we've moved into working with numpy arrays, and that the user-item matrix is very sparse.

At this point, we began implementing methods of collaborative filtering. As mentioned earlier, a first step in collaborative filtering methods for recommender systems is to determine most similar elements when given an element. In other words, if given a user, we find the most similar

users to that one. Often this similarity measure is calculated for each user (or item) against all others, resulting in a similarity matrix.

Some measures of similarity handle sparsity better than others. In our case, we don't want two rows (users) who have reviewed none of the same items to be considered similar because both have so many products unrated (0's), this wouldn't make sense. Cosine similarity is a measure that is robust against this problem, 0's essentially don't count towards the resulting similarity, not ignored but effectively neutralized.

Since we are delivering recommendations based on similar users or similar items, we need to know the similarity of users to one another and items to one another. The result is two square matrices: a user similarity matrix and item similarity matrix. Row numbers and column numbers of user similarity matrix correspond to rows of the original data (users), and rows/columns of the item similarity matrix correspond to the column indices (products). This will be useful when we need to pull actual user reviews and get info on the products.

See the code below for generating user and item similarity matrices. Note that we implement the pairwise cosine similarity function from scikit-learn. Generating the item similarity matrix is as easy as transposing the user-item matrix and doing the same calculations.

```
#Creating user similarity matrix
user_sim_mat = metrics.pairwise.cosine_similarity(nonan_test)

#Creating item similarity matrix
nonan_test_t = nonan_test.T
itm_sim_mat = metrics.pairwise.cosine_similarity(nonan_test_t)
```

It should be noted that calculating similarity matrices can be time-intensive. To save time, we tried using a smaller function that only performs element-to-element similarity. This could be used in a loop to find the most similar element (eg user) to a given input element in a recommender function. This didn't actually save much time though.

See code below for determining cosine similarity:

```
20 #function to find similarity between two users
21 def find_sim(nonan_test,u1,u2) :
22
23     cosine_sim = metrics.pairwise.cosine_similarity(nonan_test[u1], nonan_test[u2])# returns only a number
24     return cosine_sim
```

For clarity, it makes sense to consider each collaborative filtering method separately.

User-based collaborative filtering

Once familiar with the steps, the process behind user-based collaborative filtering isn't too complex. In short, we're delivering our recommendations from products that similar users liked. In other words, given a user, we find similar users and take their word for it on what products are nice.

The first step is, given a user, finding users most similar to that user. This means finding users who reviewed the same products as the input user in the same ways. Next, we find what products these similar users reviewed. We now have a list of potential recommendations: of these products reviewed by similar users, which have the best rating among those users? These will be the products recommended to our input user.

We can control how many similar users to consider and how many recommendations to give.

See code below, that finds k most similar users:

```
58 #function to find k similar users based on precalculated user similarity matrix
59 def getUserPC(user_sim_mat,u,k):
60     score=[]
61     simUsers = []
62     for others in range(len(user_sim_mat)):
63         if others != u:
64             score.append([user_sim_mat[u][others], others])
65     score.sort()
66     score.reverse()
67     for item in score[:k]:
68         simUsers.append(item[1])
69     return simUsers
70 #return score[:k]
```

The code below will use getUserPC to find most similar users, then get their reviewed products (that the input user hasn't already reviewed), and deliver recommendations:


```

72 #function to serve user based recommendations
73 def userBasedRecs(nonan_test, simUsers, unrated, numOfRecs, rated, unique_products, metadata):
74     all_simusers_productsReviews = {}
75     user_based_rec = []
76     final_rec = []
77     for user in simUsers:
78         simUsers_productreviews = {}
79         rated_items = np.where(nonan_test[user] != 0)[0]
80         rated_item_reviews = np.extract(nonan_test[user] != 0, nonan_test[user])
81         for i in list(zip(rated_items, rated_item_reviews)):
82             for item in unrated:
83                 if item == i[0]:
84                     simUsers_productreviews[i[0]] = i[1]
85             #for item in unrated:
86             # if item not in simUsers_productreviews: del simUsers_productreviews[item]
87             #print(simUsers_productreviews)
88         for key in simUsers_productreviews:
89             #If this product is not already in the master sim_users dictionary ...
90             #... create a key-value pair for that product in the master list: make sure the value is a list
91             if key not in all_simusers_productsReviews:
92                 all_simusers_productsReviews[key] = [simUsers_productreviews[key]]
93             else:
94                 #If the user's product is already in the master sim_users dictionary
95                 #Find that product in the master dict and append the user's review to the existing list of reviews for that product
96                 all_simusers_productsReviews[key].append(simUsers_productreviews[key])
97         #print(all_simusers_productsReviews)
98         for key in all_simusers_productsReviews:
99             user_based_rec.append([np.mean(all_simusers_productsReviews[key]), key])
100     user_based_rec.sort()
101     user_based_rec.reverse()
102     for item in user_based_rec[:numOfRecs]:
103         final_rec.append(item[1])
104
105     #print(user_based_rec)
106     #print(final_rec)
107
108     all5 = []
109     for item in user_based_rec:
110         if item[0] == 5.0:
111             all5.append(item[1])
112
113     for item in rated:
114         asin_pd = unique_products[item]
115         product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
116         product = str(product)
117         print("User Rated: " + product)
118     print('\n')
119     for item in final_rec:
120         asin_pd = unique_products[item]
121         product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
122         product = str(product)
123         print("Recommendation: " + product)
124
125     return final_rec

```

Let's see this in action.

Here is the output for a random user, serving 3 recommendations based on 5 most similar users

```

In [42]: # One large function: userBasedRecs
         reload(our_funcs)

         our_funcs.userBasedRecs(nonan_test, simUsers, unrated, num_rec, rated, unique_products, metadata)

```

```

User Rated: Embassy Italian Stone Design Genuine Leather Tote Bag
User Rated: Wall Earring Holder Jewelry Organizer Closet Jewelry Storage Rack - Earring Angel (Clear)
User Rated: Black Feather Boa
User Rated: UGG Australia Stain & Water Repellent, One Bottle
User Rated: Sterling Silver Cubic Zirconia Medium Round Hoop Earrings
User Rated: Leg Avenue Wet Look Leggings with Elastic Lace up

```

```

Recommendation: EVogues Apparel Women's Plus Size Glitter Necklace Accented O-Ring Strap Top
Recommendation: Timberland PRO Men's Pitboss "Steel-Toe Boot
Recommendation: Sterling Silver Multi-Color Amber Drop Earrings

```

Item-based collaborative filtering

What if, instead of getting recommendations from similar users, we took an approach of finding products similar to one we know we liked? That's the intuition behind item-based collaborative filtering.

The steps are a bit shorter than user-based. First, we get an item that a given user likes, and calculated that item's cosine similarity to all other products. After finding the most similar products, deliver ones the user hasn't already reviewed as our recommendations. We don't need to find rating information of these products since they are similar to one our user liked.

Use the item similarity matrix to find most similar items to a user's top-reviewed product, and deliver recommendations based on the n most similar. See code below:

```
127 #Item-based recommender
128 #All-in-one: most similar items --> recommended items
129 def itemSim(nonan_test, u, num_recs, unique_products, metadata, itm_sim_mat, unRate, rate):
130
131
132     #unRate, rate = getUnrated(nonan_test,u)
133     user_rev_list = []
134     simItems = []
135     #find their top rated item
136     for prod in rate:
137         user_rev = nonan_test[u,prod]
138         user_rev_list.append([user_rev,prod])
139         user_rev_list.sort()
140         user_rev_list.reverse()
141
142     most_liked_prod = user_rev_list[0][1]
143     asin_pd = unique_products[most_liked_prod]
144     product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
145     product = str(product)
146     print("Most Liked Product: " + product+ '\n')
147
148     score=[]
149     simItems = []
150     for others in range(len(itm_sim_mat)):
151         if others != most_liked_prod:
152             score.append([itm_sim_mat[most_liked_prod][others], others])
153     score.sort()
154     score.reverse()
155     #print(score[:20])
156     for item in score[:20]:
157         if item[1] in rate:
158             #print(item[1])
159             score = list(filter((item).__ne__, score))
160     #print(score)
161
162     for item in score[:num_recs]:
163         simItems.append(item[1])
164     #print(simItems)
165     for item in simItems:
166         asin_pd = unique_products[item]
167         product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
168         product = str(product)
169         print("Recommended Product: " + product)
170     #return score[:k]
```

Here is the output of the item-based recommender, serving 3 recommendations


```
In [43]: #this function takes a user input, figures out the best-reviewed item by that user ...  
#... then uses that best-reviewed-item  
  
reload(our_funcs)  
  
our_funcs.itemSim(nonan_test, u1, num_recs, unique_products, metadata)  
  
Most Liked Product: Leg Avenue Wet Look Leggings with Elastic Lace up  
  
Recommended Product: STEVEN by Steve Madden Women's Intyce Riding Boot  
Recommended Product: Dunham by New Balance Men's St. Johnsbury Sandal  
Recommended Product: 9884 Olive Drab Low Profile Death Spade Baseball Cap
```

Hybrid collaborative filtering

Often we have more context than just “who reviewed what how”. There’s information that is relevant to give personalized recommendations: information about the user or about the products: maybe a user reviews items from one brand very highly, or only buys items that have some specific feature (Kosher foods, for example). Simple review data could miss this.

In recommender systems, we can call recommender systems based on this information ‘content-based’. A common example is movie recommendations based on actors who appear in movies a user enjoyed, or based on genres the user watches most.

We used some product metadata to augment our user-based recommender system, resulting in a ‘hybrid’ of user and content-based collaborative filtering. Every product in the dataset has metadata related to product category. Naturally for our dataset, all fall under the Clothing, Shoes & Jewelry categories, but a pair of running shorts might also have categories like “Athletic wear” or “Women’s apparel”. Taking the categories for every product a user has reviewed, we can make a kind of ‘profile’ of each user. Next, after performing user-based collaborative filtering, finding products reviewed by similar users, we use the product categories of potentially recommended products and calculate the overlap between each product and the input user’s ‘profile’. In essence, we’re seeing if a product is in-line with the types of things the user buys.

For example, given a user, we find two similar users. These two users have reviewed 4 products that the user hasn’t. Two of them have an average rating of 5 stars among these 2 similar users. Our input user has bought products that have categories “Costume” “Costume jewelry” and “Clown shoes”. One of the items with 5 stars has “Costume” in common, while the other has none in common. This overlap acts as a weight on the reviews, punishing products that aren’t similar to the input user’s profile. This is a useful tiebreaker, then.

Code below:

Function for finding the category metadata of a product:

```

156 #Function that takes the col indices of items rated by a user (aka one of the outputs of getUnrated) ...
157 # .. and returns a list of unique product categories (our user's 'content profile' for hybrid method)
158 def getCate(rated, unique_products, metadata):
159     listn = []
160     for item in rated:
161         #print(item)
162         asin_pd = unique_products[item]
163         #print(asin_pd)
164         cat = metadata[metadata['asin']==asin_pd].categories.item()# 'categories']]
165         for i in cat:
166             for j in i:
167                 #print(j)
168                 listn.append(j)
169     unl = list(set(listn))
170     return unl

```

Function for determining user-based recommendations, but before delivering, weighting by the overlap of each potentially recommended product's categories w/ the input user's category 'profile'

```

173 #Hybrid recommender
174 def hybridRec(nonan_test, simUsers, unrated, numOfRecs, rated, unique_products, metadata):
175     all_simusers_productsReviews = {}
176     user_based_recs = []
177     final_recs = []
178
179     catUser = getCate(rated, unique_products, metadata)
180
181     for user in simUsers:
182         simUsers_productreviews = {}
183         rated_items = np.where(nonan_test[user] != 0)[0]
184         rated_item_reviews = np.extract(nonan_test[user] != 0, nonan_test[user])
185         for i in list(zip(rated_items, rated_item_reviews)):
186             for item in unrated:
187                 if item == i[0]:
188                     simUsers_productreviews[i[0]] = i[1]
189             #for item in unrated:
190             # if item not in simUsers_productreviews: del simUsers_productreviews[item]
191             #print(simUsers_productreviews)
192         for key in simUsers_productreviews:
193             #If this product is not already in the master sim_users dictionary ...
194             #... create a key-value pair for that product in the master list: make sure the value is a list
195             if key not in all_simusers_productsReviews:
196                 all_simusers_productsReviews[key] = [simUsers_productreviews[key]]
197             else:
198                 #If the user's product is already in the master sim_users dictionary
199                 #Find that product in the master dict and append the user's review to the existing list of reviews for that product
200                 all_simusers_productsReviews[key].append(simUsers_productreviews[key])
201         #print(all_simusers_productsReviews)
202         for key in all_simusers_productsReviews:
203             user_based_recs.append([np.mean(all_simusers_productsReviews[key]), key])
204         print("Original user-based reocs:\n", user_based_recs)
205         for item in user_based_recs:
206             #print([item[1].item()])
207             catItem = getCate([item[1].item()], unique_products, metadata)
208             match = list(set(catUser).intersection(catItem))
209             weight = len(match)/len(catItem)
210             item[0] = item[0]*weight
211
212         user_based_recs.sort()
213         user_based_recs.reverse()

```

```

214
215
216     for item in user_based_recs[:numOfRecs]:
217         final_recs.append(item[1])
218
219     print("With content-based weighting:\n", final_recs)
220
221     for item in rated:
222         asin_pd = unique_products[item]
223         product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
224         product = str(product)
225         print("User Rated:" + product)
226
227     for item in final_recs:
228         asin_pd = unique_products[item]
229         product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
230         product = str(product)
231         print("Recommendation:" + product)
232
233     return final_recs

```

Here is output of the hybrid recommender system, 3 recommendations based on 5 most similar users:

```

In [45]: #Recommendations from Hybrid Recommender
         reload(our_funcs)

         our_funcs.hybridRec(nonan_test, simUsers, unrated, num_recs, rated, unique_products, metadata)

User Rated: Embassy Italian Stone Design Genuine Leather Tote Bag
User Rated: Wall Earring Holder Jewelry Organizer Closet Jewelry Storage Rack - Earring Angel (Clear)
User Rated: Black Feather Boa
User Rated: UGG Australia Stain & Water Repellent, One Bottle
User Rated: Sterling Silver Cubic Zirconia Medium Round Hoop Earrings
User Rated: Leg Avenue Wet Look Leggings with Elastic Lace up

Recommendation: Bison Designs 38mm wide Light Duty Last Chance Belt with Gunmetal Buckle
Recommendation: Casio Men's AW49HE-2AV Ana-Digi Dual Time Watch
Recommendation: Sterling Silver Multi-Color Amber Drop Earrings

```

Evaluation

So we have delivered recommendations and we can kind of intuitively tell if they make sense or not, but what are standard, quantifiable ways of evaluating a recommender system?

Three common methods are mean absolute error, precision, and recall. MAE is an evaluator of predictions (continuous) while precision and recall evaluate classifications (categorical), just as a point of distinction. We used MAE and recall to evaluate all three recommender systems that we designed.

Mean absolute error

Mean absolute error in the context of recommender systems means predicting how a user will review a product, and evaluating how far off our prediction was. Doing this many times and getting an average is the 'mean' in mean absolute error.

We find a product the input user has reviewed, withhold it as if it were unrated by the user, and depending on the method, find a way of predicting its rating.

For user-based, this involves finding the most similar users who have reviewed the withheld product, and getting their average rating. For item-based, this is a little more convoluted. We withhold one of the items that a given user has reviewed, then find the k most similar items *among the products reviewed by the user*. The user's rating average of these products is what we would predict is given to the withheld product. MAE for the hybrid method is essentially the same as user-based MAE, except that we apply the content-based weighting to both the withheld item and the average reviews of the item from most similar users. In essence this just changes the scale, but the trends on error would be the same for both methods as you tweaked number of similar users.

See below, each method has both a base code for calculating absolute error between actual review and predicting review, and a function for performing this across a subset of the data and taking an average:

User-based

```
453 #Given a user, find the most similar users who have also reviewed that user's 1st reviewed product ...
454 #Get the avg rating of that product from the k most similar users, that's our prediction
455 #Returns avg rating and actual rating ... this function is used for calculating MAE
456 def collabo_evaluator(user, k, nonan_test):
457     simuser = []
458     simu = []
459     users_products = list(np.nonzero(nonan_test[user]))
460     the_prod = users_products[0][0] #arbitrary product from their list of products
461     #print(the_prod)
462     users_review = nonan_test[user, the_prod]
463     #print(users_review)
464     #find all users who have reviews 'the_prod'
465     users_with_prod = np.where(nonan_test[:, the_prod] != 0)[0]
466     for i in users_with_prod:
467         if i != user:
468             sim = find_sim(nonan_test, user, i)
469             simuser.append([sim, i])
470     simuser.sort()
471     simuser.reverse()
472     for item in simuser[:k]:
473         simu.append(item[1])
474     #print(simu)
475     rat = 0
476     for kuser in simu:
477         rat += nonan_test[kuser][the_prod]
478     pred_rating = rat/k
479     #print(pred_rating)
480     return pred_rating, users_review
```

Get MAE for user-based collaborative filtering on a subset of the data

```
561 #performs many hybrid runs, using collabo_evaluator, returns mae
562 def calc_mae_user(nonan_test, ratio, k):
563     #Get the subset of users
564     num_users = np.shape(nonan_test)[0]
565     list_of_users = list(range(0, num_users))
566     np.random.shuffle(list_of_users)
567     subset_size = int(num_users * ratio)
568     subset_users = list_of_users[0:subset_size]
569     #print(subset_users)
570     abserror = 0
571     count = 0
572     for wuser in subset_users:
573         pred, userr = collabo_evaluator(wuser, k, nonan_test)
574         abserror += abs(pred-userr)
575         #print(abserror)
576         count += 1
577     MAE = abserror/count
578     print("MEAN ABSOLUTE ERROR: ", MAE)
```

Here's the results when evaluation user-based recommender (5 most similar users) on 0.1% of the data:

```
#User-based version of MAE

reload(our_funcs)

our_funcs.calc_mae_user(nonan_test, 0.01, n)

MEAN ABSOLUTE ERROR: 0.796946564885
```


Item-based

Calculating absolute error

```
448 #MAE for item-based
449 #We need to:
450 #Get a user's product reviews
451 #Pop one of their products and find most similar products to it (that the user has reviewed!)
452 #Get the average rating that the user has given to the most similar k items
453
454 def itemRec_evaluator(user, k, nonan_test, itm_sim_mat):
455
456     score = []
457     simuser = []
458     simu = []
459     simItems = []
460     users_products = list(np.nonzero(nonan_test[user]))
461     the_prod = users_products[0][0] #arbitrary product from their list of products - popped product
462     #print(the_prod)
463
464     users_review = nonan_test[user, the_prod]
465     #print(users_review)
466
467     #find the products rated by user and its similarity to popped product
468     unRate, rate = getUnrated(nonan_test, user)
469     #print(rate)
470     for item in rate:
471         if item != the_prod:
472             score.append([itm_sim_mat[the_prod][item], item])
473     score.sort()
474     score.reverse()
475     #print(score, '\n')
476     for item in score[:k]:
477         simItems.append(item[1])
478
479     itmRatings = 0
480
481     for item in simItems:
482         itmRatings += nonan_test[user][item]
483     pred_rating = itmRatings/k
484     #print(pred_rating)
485     #print(users_review)
486     return pred_rating, users_review
```

Calculate MAE by performing the above function over a subset

```
488 def calc_mae_itm(nonan_test, ratio, k):
489
490     #Get the subset of users
491     num_users = np.shape(nonan_test)[0]
492     list_of_users = list(range(0, num_users))
493     np.random.shuffle(list_of_users)
494     subset_size = int(num_users * ratio)
495     subset_users = list_of_users[0:subset_size]
496     #print(subset_users)
497     abserror = 0
498     count = 0
499     for wuser in subset_users:
500         pred, userr = itemRec_evaluator(wuser, k, nonan_test)
501         abserror += abs(pred-userr)
502         #print(abserror)
503         count += 1
504     MAE = abserror/count
505     print("MEAN ABSOLUTE ERROR: ", MAE)
```


Hybrid

Calculating absolute error

```
483 def collabo_evaluator_hyb(user, k, nonan_test, unique_products, metadata):
484
485     simuser = []
486     simu = []
487     rated_prod = np.where(nonan_test[user,:] != 0)[0]
488     #print(rated_prod)
489     users_products = list(np.nonzero(nonan_test[user])) #get products rated by user
490     #print(users_products)
491     #categories of products rated by user expect the product we are predicting for
492     catUser = getCate(rated_prod[1:], unique_products, metadata)
493
494     the_prod = users_products[0][0] #arbitrary product from their list of products
495     users_review = nonan_test[user, the_prod] #Users review for that product
496
497     catItem = getCate([the_prod], unique_products, metadata)
498     match = list(set(catUser).intersection(catItem))
499     weight = len(match)/len(catItem)
500     #print(weight)
501     users_review = users_review*weight #So we have weighted user review based on how similar that item is to all rated items
502     #print(the_prod)
503
504     #print(users_review)
505     #find all users who have reviews 'the_prod'
506     users_with_prod = np.where(nonan_test[:, the_prod] != 0)[0]
507     for i in users_with_prod:
508         if i != user:
509             sim = find_sim(nonan_test,user,i)
510             simuser.append([sim, i])
511     simuser.sort()
512     simuser.reverse()
513
514     for item in simuser[:k]:
515         simu.append(item[1])
516
517     #print(simu)
518
519     rat = 0
520     othusers_review = 0
521     #For every similar user, get their rating, and determine the average rating
522     for kuser in simu:
523         rat = nonan_test[kuser][the_prod]
524         othrated_prod = np.where(nonan_test[kuser,:] != 0)[0]
525         othrated_prod = list(filter((the_prod).__ne__, othrated_prod))
526         othusers_products = list(np.nonzero(nonan_test[kuser])) #list of products rated by similar user
527         othusers_cat = getCate(othrated_prod, unique_products, metadata)
528         match = list(set(othusers_cat).intersection(catItem))
529
530         #print(match)
531         #print(othusers_cat)
532         #print(catItem)
533         weight = len(match)/len(catItem)
534         #print(weight)
535         othusers_review += rat*weight
536
537     pred_rating = othusers_review/k
538     #print(pred_rating)
539     #print(users_review)
540     return pred_rating, users_review
```

Calculate MAE for hybrid method

```
541 #performs many hybrid runs, using collabo_evaluator, returns mae
542 def calc_mae_hyb(nonan_test, ratio, k, unique_products, metadata):
543
544     #Get the subset of users
545     num_users = np.shape(nonan_test)[0]
546     list_of_users = list(range(0,num_users))
547     np.random.shuffle(list_of_users)
548     subset_size = int(num_users * ratio)
549     subset_users = list_of_users[0:subset_size]
550     #print(subset_users)
551     abserror = 0
552     count = 0
553     for wuser in subset_users:
554         pred, userr = collabo_evaluator_hyb(wuser, k, nonan_test, unique_products, metadata)
555         abserror += abs(pred-userr)
556         #print(abserror)
557         count += 1
558     MAE = abserror/count
559     print("MEAN ABSOLUTE ERROR: ", MAE)
```

Here are results when evaluating hybrid recommender (5 most similar users) on 1% of the data:

```
#Hybrid version of MAE
reload(our_funcs)

our_funcs.calc_mae_hyb(nonan_test, 0.01, n, unique_products, metadata)

MEAN ABSOLUTE ERROR: 0.893930322129
```

Results

MAE, on 20% of the data (5% for hybrid method), adjusting parameters for number of similar users and number of products to recommend.

Note that this is taking a random subset of the data, so scores mildly fluctuate.

n: number of similar users/products to base recommendations off of

n=2	Method	ratio = 0.2 (0.05 for hybrid)
	User	0.7554
	Item	0.6938
	Hybrid	0.82168

n=5	Method	ratio = 0.2 (0.05 for hybrid)
	User	0.8859
	Item	0.7319
	Hybrid	0.8938

n=8	Method	ratio = 0.2 (0.05 for hybrid)
	User	1.1118
	Item	1.344
	Hybrid	1.08466

We see that MAE increases as we take more neighbors (either similar users or similar products). This makes sense, as we loosen the restriction on whose reviews can be considered when predicting a given user's review of a product, you'd think the less similar users would throw off the accuracy of prediction. One way of alleviating this, and improving the performance of the recommender systems in general, would be the weight the 'voice' a user is given in recommendations based on their similarity. The best performance is on item-based recommendation when the number of similar items is kept to a minimum.

Recall

Recall is about testing whether or not our recommendations are relevant. We hold out a product or several products that a given user liked, and we see if these products show up in our recommendations. Normally for our recommender systems, we wouldn't want to be serving up products the user has already purchased, so some additional work needs to be done to handle this.

Recall is then: ($\#$ of withheld products in recommendations) / ($\#$ of withheld products) . Because users have only reviewed a small number of products, we only withheld one product, and then repeated the process over a subset of the data to get an average recall for each recommender system.

Understanding how recall works, it becomes clear that there's a way to "game" the measure: if you serve up a large amount of recommendations, you are more likely to find the holdout

product(s). That's worth keeping in mind: how the parameters will affect the measure, and what our priorities are (optimizing recall? Delivering a reasonable number of recommendations?).

See code below:

User-based

```
307 def user_based_recall(nonan_test, ratio, numSimUsers, numRecs, unique_products, metadata, user_sim_mat):
308     #get a subset of users
309     num_users = np.shape(nonan_test)[0]
310     list_of_users = list(range(0,num_users))
311     np.random.shuffle(list_of_users)
312     subset_size = int(num_users * ratio)
313     subset_users = list_of_users[0:subset_size]
314
315     user = 0
316     user_recall = 0
317
318     #for each index in that subset_users, pull the user at that row ...
319     # ... and do process of getting similar users
320     for wuser in subset_users:
321         user+=1
322         simUsers = getUserPC(user_sim_mat, wuser, numSimUsers)
323         unrated, rated = getUnrated(nonan_test, wuser)
324         wuser_reviews_list = []
325         for product in rated:
326             wuser_review = nonan_test[wuser, product]
327             wuser_reviews_list.append([wuser_review, product])
328         wuser_reviews_list.sort()
329         wuser_reviews_list.reverse()
330
331         #We now have the users product reviews and similar users.
332         #We want to add validation prods to 'unrated' list, and serve recommendations to user
333         validation_prods = wuser_reviews_list[0][1]
334         #print(validation_prods)
335         unrated.append(validation_prods)
336         #print(unrated)
337         recommendations = userBasedRecs(nonan_test, simUsers, unrated, numRecs, rated, unique_products, metadata)
338         recommendations = map(int,recommendations)
339         validation_prods = int(validation_prods)
340         count = 0
341         for rec in recommendations:
342             if rec == validation_prods:
343                 count+=1
344         user_recall += count/1
345         #print(count)
346         #print(user_recall)
347         #print(user)
348
349     final_recall = user_recall/user
350
351     print(final_recall)
```

Item-based

```
353 def item_based_recall(nonan_test, u, ratio, numSimUsers, numRecs, unique_products, metadata, itm_sim_mat):
354
355     num_users = np.shape(nonan_test)[0]
356     list_of_users = list(range(0, num_users))
357     np.random.shuffle(list_of_users)
358     subset_size = int(num_users * ratio)
359     subset_users = list_of_users[0:subset_size]
360
361     user = 0
362     user_recall = 0
363     for wuser in subset_users:
364         user += 1
365         unRate, rate = getUnrated(nonan_test, wuser)
366         user_rev_list = []
367
368         #find their top rated item
369         for prod in rate:
370             user_rev = nonan_test[wuser, prod]
371             user_rev_list.append((user_rev, prod))
372             user_rev_list.sort()
373             user_rev_list.reverse()
374         print(user_rev_list)
375         if len(user_rev_list) < 2:
376             print("Can't perform recall!")
377         else:
378             most_liked_prod = user_rev_list[0][1]
379             #validation_prod = user_rev_list[0][1]
380             unRate.append(most_liked_prod)
381             asin_pd = unique_products[most_liked_prod]
382             product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
383             product = str(product)
384             print("Most Liked Product:" + product)
385
386             score=[]
387             simItems = []
388
389             for item in rate:
390                 if item != most_liked_prod:
391                     #score.append((find_sim(nonan_test.T, most_liked_prod, item), item))
392                     score.append((itm_sim_mat[most_liked_prod][item], item))
393             score.sort()
394             score.reverse()
395
396             rate = []
397
398             for item in score:
399                 rate.append(item[1])
400             print(rate)
401
402
403             simItems = []
404
405             prod_to_use = score[0][1]
406
407             print(prod_to_use)
408
409             score=[]
410
411             for others in range(len(itm_sim_mat)):
412                 if others != prod_to_use:
413                     score.append((itm_sim_mat[prod_to_use][others], others))
414             score.sort()
415             score.reverse()
416
417             print(score[:20])
418             for item in score[:20]:
419                 if item[1] in rate:
420                     score = list(filter((item).__ne__, score))
421             #print(score)
422
423             for item in score[:numRecs]:
424                 simItems.append(item[1])
425             #print(simItems)
426             '''
427             for item in simItems:
428                 asin_pd = unique_products[item]
429                 product = metadata[metadata['asin']==asin_pd].title.item()# 'categories']]
430                 product = str(product)
431                 print("Recommended Product: " + product)
432             '''
433             recommendations = map(int, simItems)
434             validation_prods = int(most_liked_prod)
435             count = 0
436             for rec in recommendations:
437                 if rec == validation_prods:
438                     count += 1
439             user_recall += count/1
440             print(count)
441             print(user_recall)
442             print(user)
443
444     final_recall = user_recall/user
445
446     print(final_recall)
```


Hybrid

```
258 #Hybrid based Recall
259 def hybrid_recall(nonan_test, ratio, numSimUsers, numRecs, unique_products, metadata, user_sim_mat):
260     #get a subset of users
261     num_users = np.shape(nonan_test)[0]
262     list_of_users = list(range(0,num_users))
263     np.random.shuffle(list_of_users)
264     subset_size = int(num_users * ratio)
265     subset_users = list_of_users[0:subset_size]
266
267     user = 0
268     user_recall = 0
269
270     #for each index in that subset_users, pull the user at that row, and do process of getting similar users
271     for wuser in subset_users:
272         user+=1
273         #simUsers = getUser(nonan_test, wuser, numSimUsers, similarity=find_sim)
274         simUsers = getUserPC(user_sim_mat,wuser,numSimUsers)
275         unrated, rated = getUnrated(nonan_test, wuser)
276         wuser_reviews_list = []
277         print("Items that user rated: ", rated)
278         for product in rated:
279             wuser_review = nonan_test[wuser, product]
280             wuser_reviews_list.append([wuser_review, product])
281         wuser_reviews_list.sort()
282         wuser_reviews_list.reverse()
283
284         #We now have the users product reviews and similar users.
285         #We want to add validation prods to 'unrated' list, and serve recommendations to user
286         validation_prods = wuser_reviews_list[0][1]
287         print("Product pulled: ", validation_prods)
288         unrated.append(validation_prods)
289         #print(unrated)
290         recommendations = hybridRec(nonan_test, simUsers, unrated, numRecs, rated, unique_products, metadata)
291         recommendations = map(int,recommendations)
292         validation_prods = int(validation_prods)
293         count = 0
294         for rec in recommendations:
295             if rec == validation_prods:
296                 count+=1
297         user_recall += count/1
298         print(count)
299         print(user_recall)
300         print(user)
301
302     final_recall = user_recall/user
303
304     print("Final recall: ", final_recall)
```


Results

Recall, on 20% of the data (only 5% for hybrid method), adjusting parameters for number of similar users and number of recommendations

(note: Item-based recommendation is not affected by the number of similar users - in this implementation: number of recs would be synonymous with number of similar products to select. Variation is due to random sampling)

n: number of similar users, ratio: subset of the data to test on, num_recs: number of recommendations to deliver

n=2	ratio = 0.2 (0.05 for hybrid)	num_recs = 2	num_recs = 5	num_recs = 8 (ratio = 0.1)
	User	0.24423	0.37557	0.41705
	Item	0.25403	0.44527	0.57085
	Hybrid	0.31566	0.410138	0.41935

n=5	ratio = 0.2 (0.05 for hybrid)	num_recs = 2	num_recs = 5	num_recs = 8
	User	0.17289	0.34792	0.44124
	Item	0.24596	0.45506	0.58352
	Hybrid	0.28341	0.44009	0.49308

n=8	ratio = 0.2 (0.05 for hybrid)	num_recs = 2	num_recs = 5	num_recs = 8
	User	0.13594	0.26958	0.3899
	Item	0.25057	0.43894	0.57142
	Hybrid	0.20737	0.37557	0.45852

Analysis of results

From the results, we see the way recall can be “gamed”: we’ll be more likely to serve the held-out validation product as a recommendation as we increase the number of recommendations. Unfortunately, the recall here never even breaks the 50% threshold: more often than not, sometimes as often as ~70% of the time, the validation product is not in the recommendations.

On the other hand, this isn’t necessarily a great measure for our data: the dataset is very sparse, and the products are diverse. Especially for item-based recommendation, where hold out an item, then perform item-based recommendation on the product *most similar to the one held out*, how often will this methodology yield the results that recall is looking for? People buy a

diverse range of clothing, shoes, and jewelry, unlike movies where we are at least confined to one format of products. If our recall function pulls out the only pair of shoes a user has reviewed, its likely that the most similar reviewed item, the one used to deliver recommendations, won't turn up that hold out product.

Conclusions

Our recommender systems were relatively successful in delivering relevant products to users. Given the sparsity of user data and the variety of products within the dataset, this encouraging. Mean absolute error was within a tolerable range for the most part (about 1-star off of the user's actual review), but when evaluated by recall the recommender systems are unimpressive. Whether or not the data even provides the opportunity for recall to be high (due to data sparsity), is an open question, but at the same time an "eyeball test" of how well the recommender systems are performing will not work at scale.

Two issues we encountered in implementing the three recommender systems and two evaluation methods were runtime and uncertainty on the best ways to leverage the metadata content provided with the data.

Runtime was especially an issue with the hybrid method. After reducing the metadata down to only those products which were in the reviews dataset, querying the metadata to retrieve product categories sped up greatly. Still, when running evaluation methods that have the recommender system performed hundreds of times, hybrid method was still very slow.

In the future, it would be interesting not only to leverage more of the metadata to make better content-based recommendations, but also to try an entirely different method of recommender systems from collaborative filter. For example, we did not explore matrix factorization in this project, and it would be worthwhile to see if that method brings any improvements to speed or accuracy/relevance of recommendations.